

CES-27 Processamento Distribuído

MapReduce

Prof Juliana Bezerra
Prof Celso Hirata
Prof Vitor Curtis

Outline

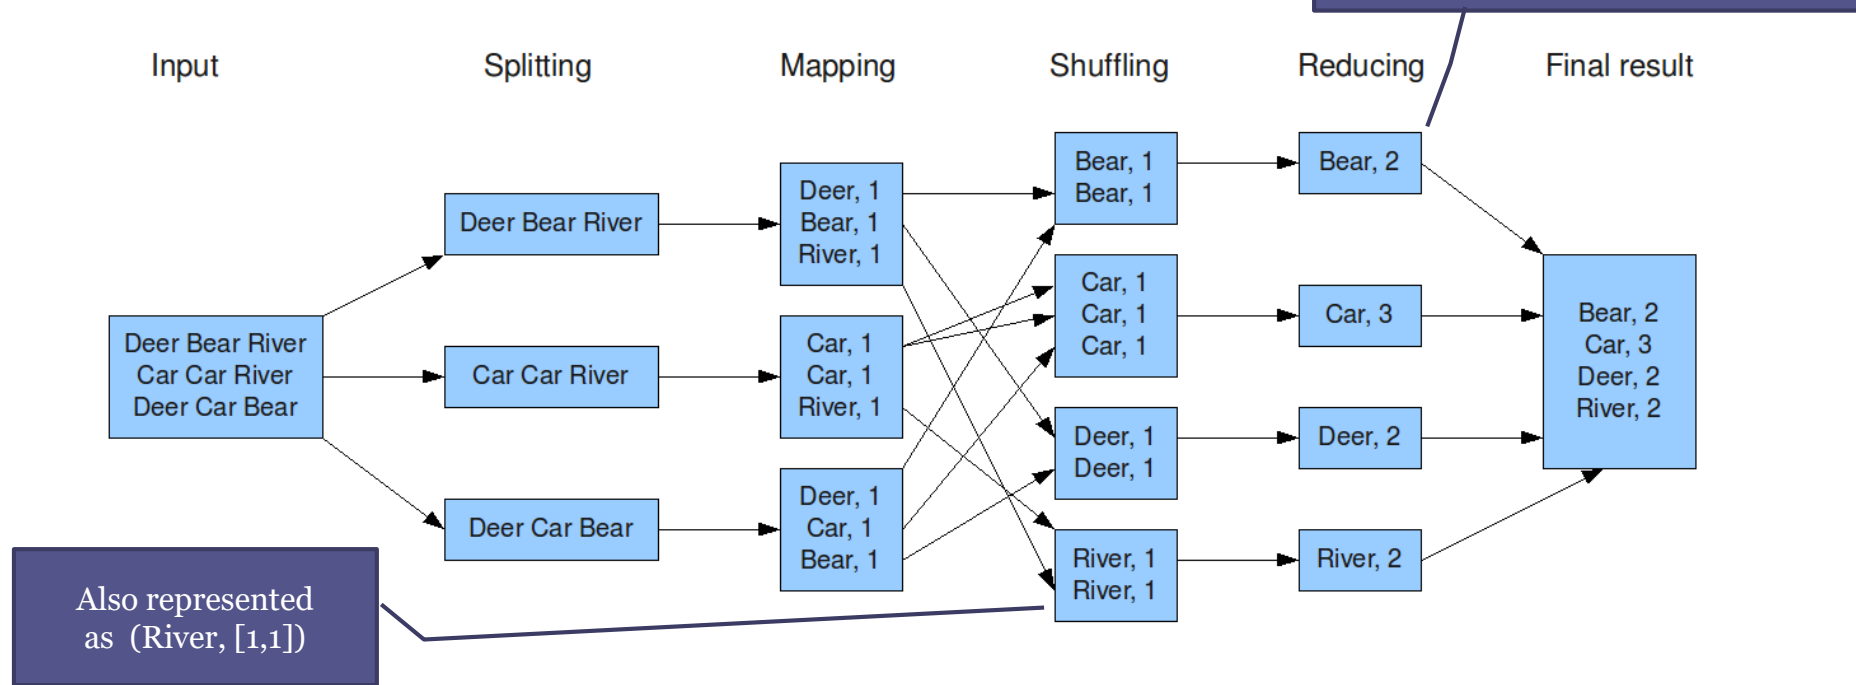
- What is MapReduce?
- MapReduce Execution Overview
- Curiosities
 - Failure Tolerance
 - Task Granularity
 - Combiner Function
- Patterns and applications

MapReduce

- Reference paper: J. Dean and S. Ghemawat, “MapReduce: Simplified Data Processing on Large Clusters”, Google, 2004
- MapReduce is a **programming model** and an associated **implementation** for processing and generating **large data sets**
 - Inspired by the map and reduce primitives present in Lisp and many other functional languages
 - Map: (input shard) → intermediate(key/value pairs)
 - A **map** operation to each logical “record” in our input in order to compute a set of intermediate **key/value pairs**
 - Reduce: intermediate(key/value pairs) → result files
 - A **reduce** operation to all the values that shared the **same key**, in order to **combine the derived data** appropriately
 - Focus: high performance
 - ... but hides the details of parallelization, fault-tolerance, data distribution and load balancing in a library

Word Count - A Typical Example

- Input: Large number of text documents (or a big one document)
- Operation: Compute word count across all the documents
- Solution
 - Mapper: For every word in a document, output (word, "1")
 - Reducer: Sum all occurrences of words and output (word, totalCount)
- A typical phase in the middle:
 - Shuffle (or partition)
 - Also referenced together to grouping/sorting

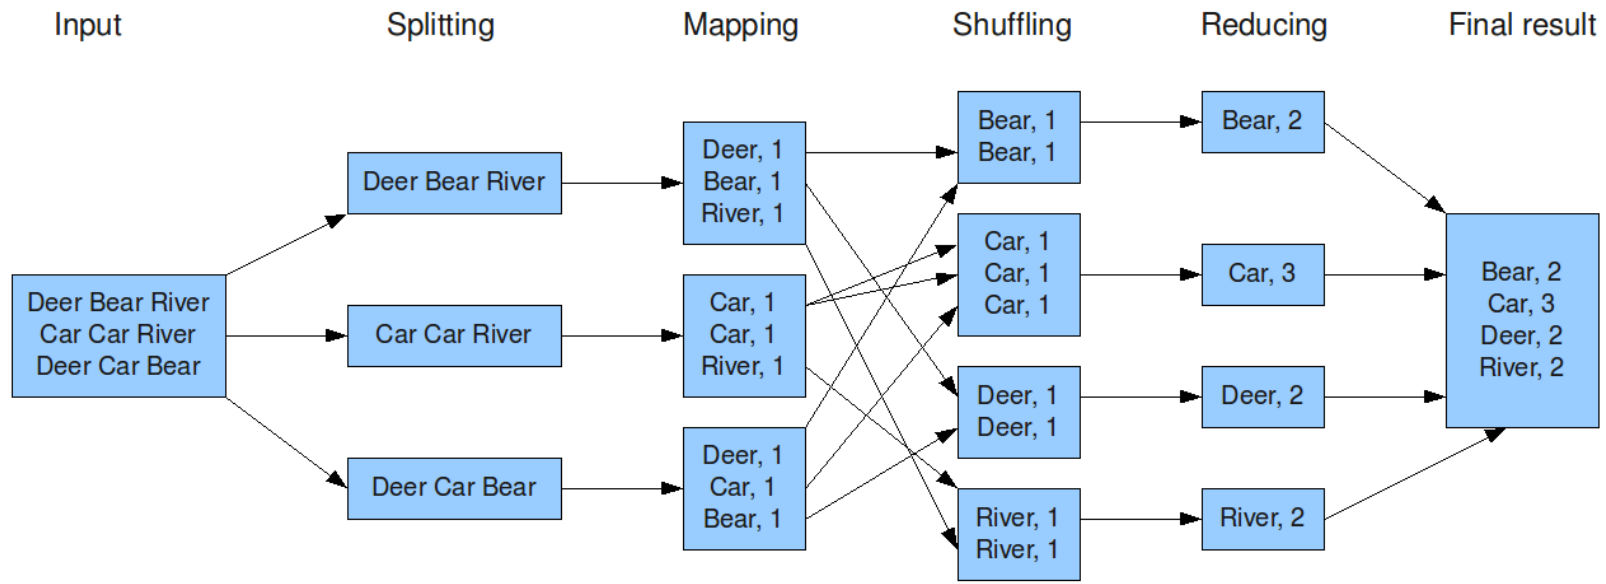


Word Count - A Typical Example

```
map(String input_key, String input_value):  
  // input_key: document name  
  // input_value: document contents  
  for each word w in input_value:  
    EmitIntermediate(w, "1");
```

```
reduce(String output_key, Iterator intermediate_values):  
  // output_key: a word  
  // intermediate values: list of counts  
  int result = 0;  
  for each v in intermediate_values:  
    result += ParseInt(v);  
  Emit(AsString(result));
```

Only v of the given
output_key



Who has MapReduce?


- Google
 - Original proprietary implementation
- Apache Hadoop MapReduce
 - Most common (open-source) implementation
 - Built based on specs defined by Google
 - Proprietary solutions:
 - Amazon Elastic MapReduce (run on Amazon EC2)
 - IBM Hadoop
- Apache Spark
 - A fast and general open-source engine for large-scale data processing
 - It has MapReduce and other things
 - SPARK SQL, SPARK Streaming, MLlib (Machine Learning) and GraphX (graph processing)
 - Spark is capable to run programs up to 100x faster than Hadoop MapReduce in memory, or 10x faster on disk

Outline

- What is MapReduce?
- MapReduce Execution Overview
- Curiosities
 - Failure Tolerance
 - Task Granularity
 - Combiner Function
- Patterns and applications

MapReduce Execution Overview


- Keep in mind:
 - map task \neq map worker \neq user-defined *Map* function
 - reduce task \neq reduce worker \neq user-defined *Reduce* function



Note: task == job

MapReduce Execution Overview - Step 1


- The MapReduce library in the user program first **splits the input files into M pieces** of typically 16 megabytes to 64 megabytes (MB) per piece (controllable by the user via an optional parameter).
- It then starts up many copies of the program on a cluster of machines \Rightarrow **fork processes**



Note: We say split or shard


MapReduce Execution Overview - Step 2

- One of the copies of the program is special – the **master** (leader).
- The rest are **workers** that are assigned work by the master.
- There are **M map tasks** and **R reduce tasks** to assign.
- The **master** picks idle workers and **assigns** each one a map task or a reduce task.



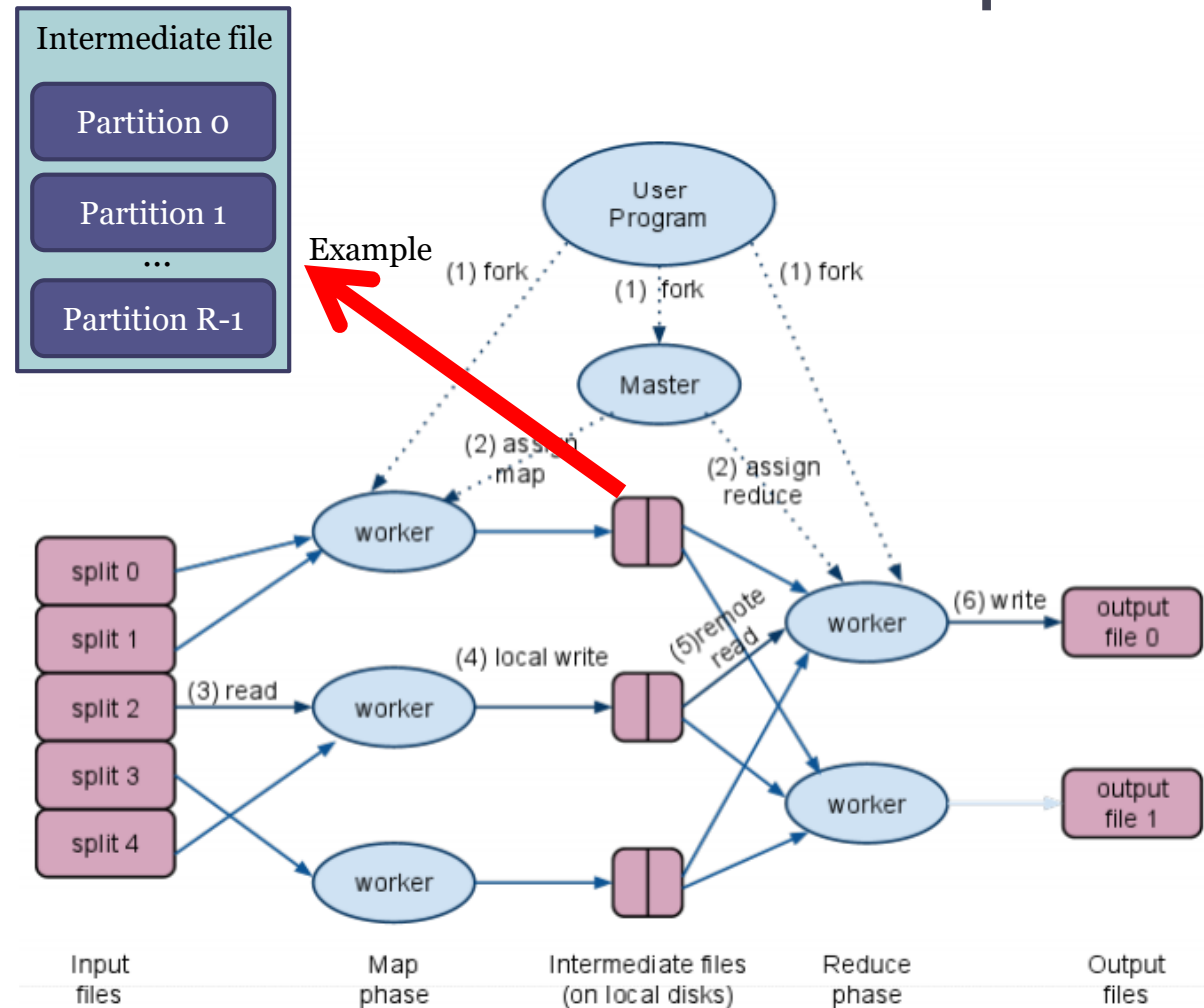
MapReduce Execution Overview - Step 3

- A worker who is assigned a **map task** reads the contents of the corresponding input split.
- It calls user-defined **Map function**.
- The **intermediate key/value pairs** produced by the *Map* function are buffered in **memory**.




MapReduce Execution Overview - Step 4

- Periodically, the buffered pairs are written to **local disk**, partitioned into **R regions** by the **partitioning function**.
 - Default function: $\text{hash}(\text{key}) \bmod R$
 - “Map worker” partitions the data by keys
 - It also decides which of R reduce workers will work on which key
- The locations of these buffered pairs on the local disk are passed back to the **master**, who is responsible for forwarding these locations to the reduce workers.




MapReduce Execution Overview - Step 5

- When a **reduce worker** is notified by the master about these locations, it uses **remote procedure calls** to **read** the buffered data from the **local disks** of the map workers.
- When a reduce worker has read all intermediate data
 - It **sorts it by the intermediate keys**
 - All occurrences of the **same key are grouped together**.
- The sorting is needed because typically many **different keys map to the same reduce task**.
 - If the amount of intermediate data is too large to fit in memory, an external sort is used.



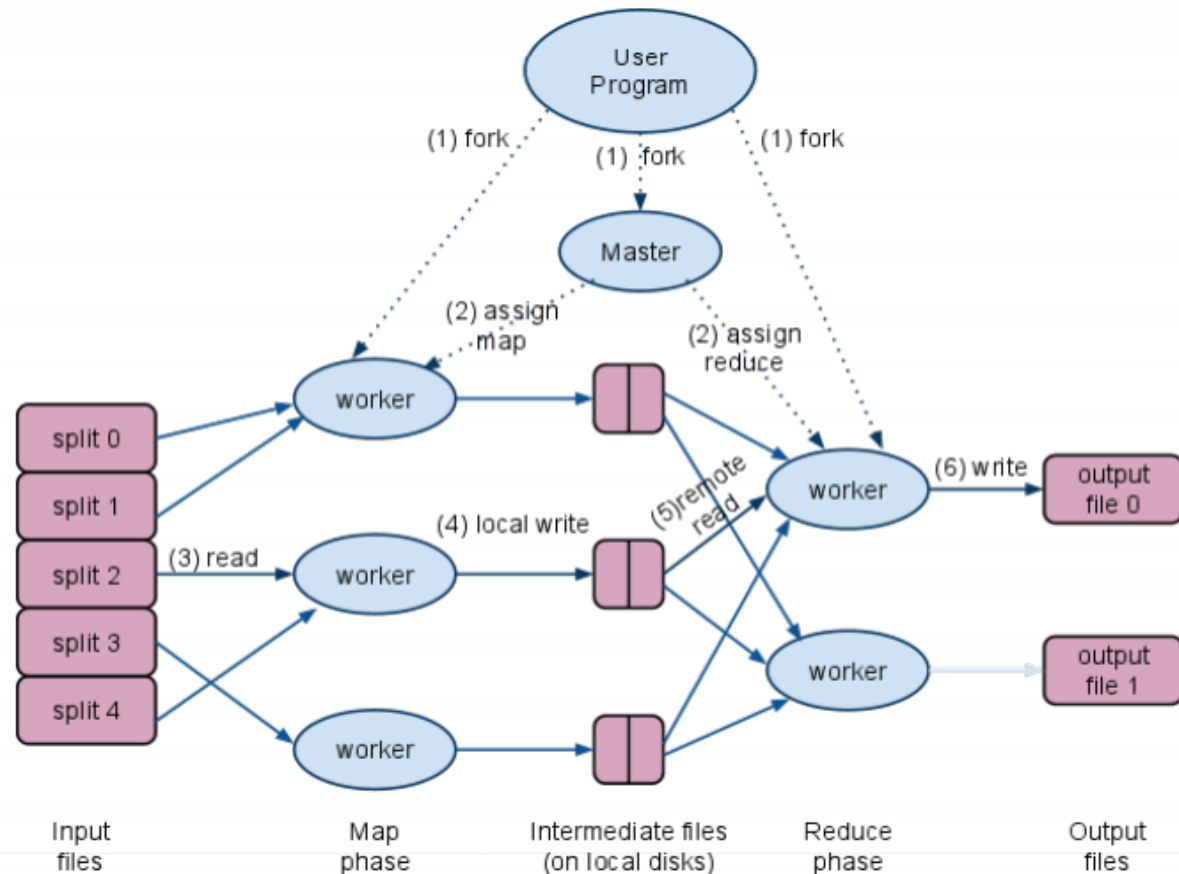
MapReduce Execution Overview - Step 6

- The reduce worker iterates over the sorted intermediate data and, for each unique intermediate key encountered,...
 - It passes the **key and the corresponding set of intermediate values** to the user's **Reduce function**.
 - $\langle \text{key}, (\text{value1}, \text{value2}, \text{value3}, \text{value4}, \dots) \rangle$
- The output of the *Reduce* function is appended to a final **output file** for this reduce partition



MapReduce Execution Overview - Step 7

- When all map tasks and reduce tasks have been completed, the **master wakes up the user program**.
 - At this point, the MapReduce call in the user program returns back to the user code.
- The output of the MapReduce execution is available in the **R output files**
 - Typically, **users do not need to combine these R output files into one file**
 - They often pass these files as input to another MapReduce call, or
 - They use them from another distributed application that is able to deal with input that is partitioned into multiple files




Outline

- What is MapReduce?
- MapReduce Execution Overview
- Curiosities
 - Failure Tolerance
 - Task Granularity
 - Combiner Function
- Patterns and applications

Failure Tolerance


- Master Data Structures
 - **State** (*idle, in-progress, or completed*) of each **map task** or **reduce task**
 - **Identity** of the **worker** machine (for non-idle tasks)
 - **Locations** and **sizes** of the **R intermediate file** regions produced by the map task
- Master Failure
 - One approach
 - Master writes periodic **checkpoints** of the **master data structures**
 - If the master task dies, a **new copy can be started** from the last **checkpointed** state.
 - Current approach (with one master)
 - To **abort** the MapReduce computation if the master fails.
 - **Clients** can check for this condition and **retry** the MapReduce operation if they desire.



Failure Tolerance

- Worker Failure


- The master **pings** every **worker** periodically.
- If **no response** in a certain amount of time, the master marks the **worker** as **failed**.



- Any map task or reduce task in progress on a failed worker is also **reset to idle** and becomes eligible for **rescheduling**.
- Any **map tasks** completed by the worker are **reset** back to their initial **idle** state, and therefore become **eligible for scheduling** on other workers.
 - Because the output of map tasks is stored on the local disk(s) of the failed machine and is therefore inaccessible.
 - Completed reduce tasks do not need to be re-executed since their output is stored in a global file system
- When a map task is executed first by worker A and then later executed by worker B (because A has failed), all workers executing reduce tasks are notified of the re-execution. Any reduce task that has not already read the data from worker A will read the data from worker B.

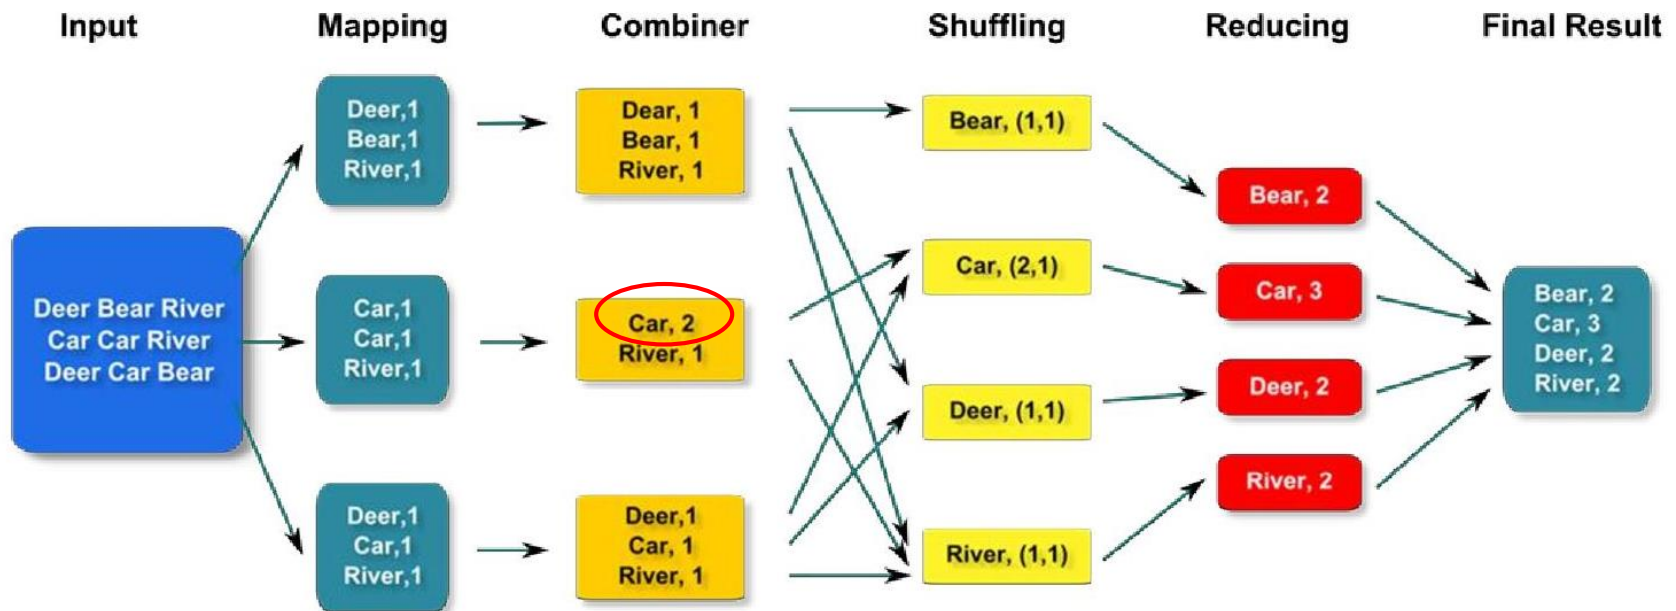
Task Granularity

- MapReduce subdivide:
 - map phase into M pieces
 - reduce phase into R pieces
- Master must:
 - Make $O(M + R)$ scheduling decisions
 - Keep $O(M * R)$ state in memory
 - However, it is **small**!
 - Approximately **one byte** of data per “map task / reduce task” pair.
- In practice:
 - Choose M so that each individual task is roughly 16 MB to 64 MB of input data
 - Make R a small multiple of the number of worker machines we expect to use.
 - $M = 200.000$ and $R = 5.000$, using 2.000 worker machines



Combiner Function


- In some cases, there is a significant **repetition in the intermediate keys** produced by each map task
 - In word count example: <the, 1>
- The **Combiner function** is executed on each machine that performs a map task
 - Typically the **same code** is used to implement both the **combiner** and the **reduce** functions



Outline

- What is MapReduce?
- MapReduce Execution Overview
- Curiosities
 - Failure Tolerance
 - Task Granularity
 - Combiner Function
- Patterns and applications

MapReduce Design Patterns



O'REILLY*

Donald Miner & Adam Stoeck

Patterns and applications

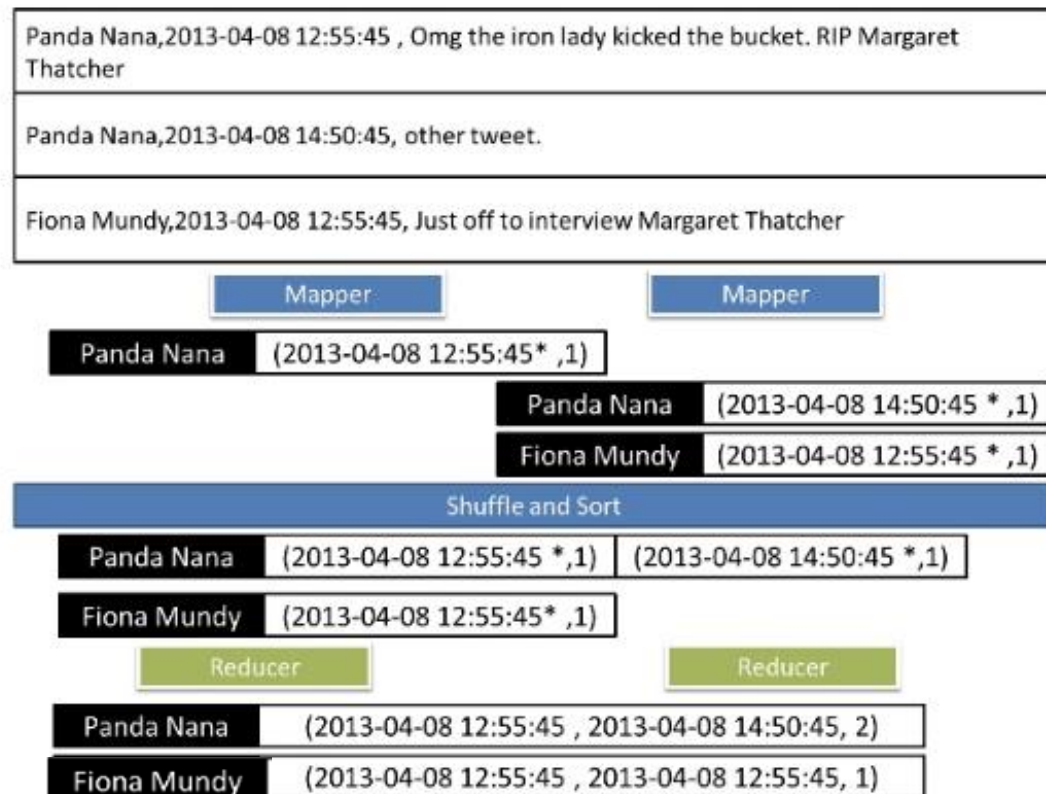
- 23 patterns grouped into six categories
 - **Summarization (*)**
 - Top-down summaries to get a top-level view
 - **Filtering (*)**
 - Extract interesting subsets of the data
 - **Data Organization**
 - Reorganize and restructure data to work with other systems or to make MapReduce analysis easier
 - **Joins**
 - Bringing and analyze different data sets together to discover interesting relationships
 - **Metapatterns**
 - Piece together several patterns to solve a complex problem or to perform several analytics in the same job
 - **Input and output**
 - Custom the way to use Hadoop to input and output data.

Numerical Summarization

- A general pattern for calculating aggregated statistical values over your data
 - To deal with numerical data or counting
 - To group data by specific fields

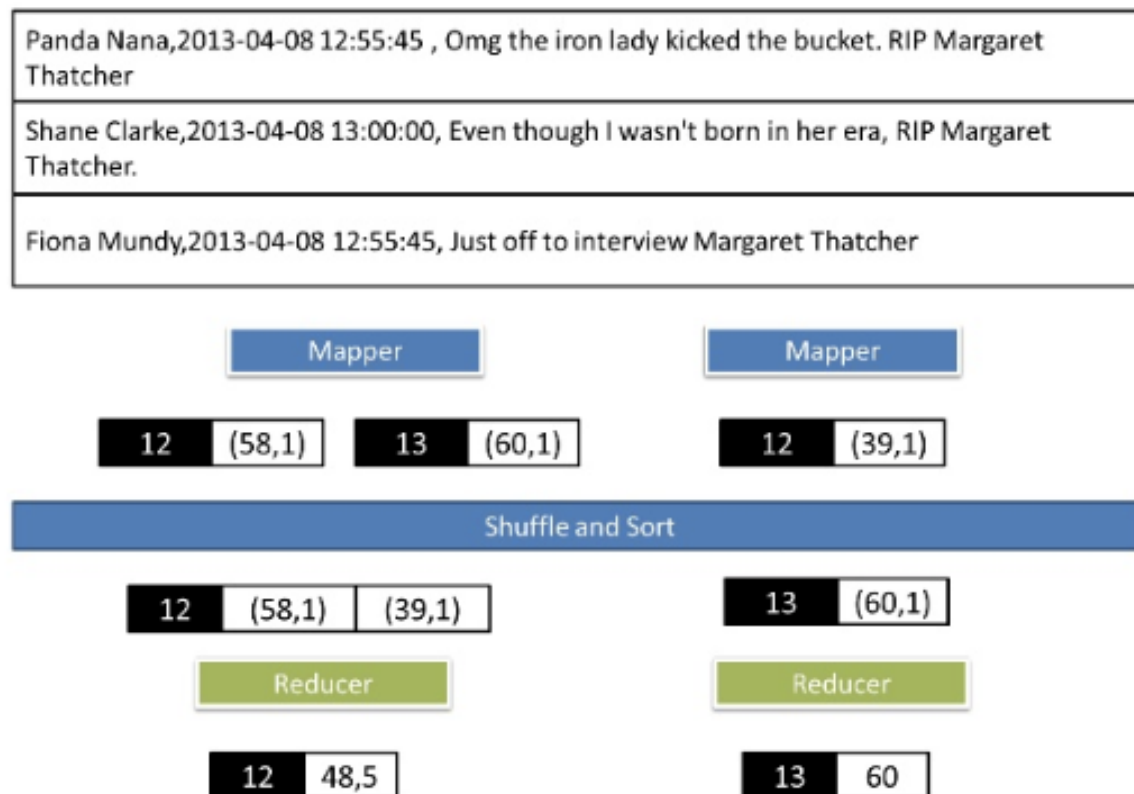
Numerical Summarization

- Word count, Record count
- Min, max, count** of a particular event
 - E.g. Given a list of tweets (username, date, text), determine first and last time an user commented and the number of times



Numerical Summarization

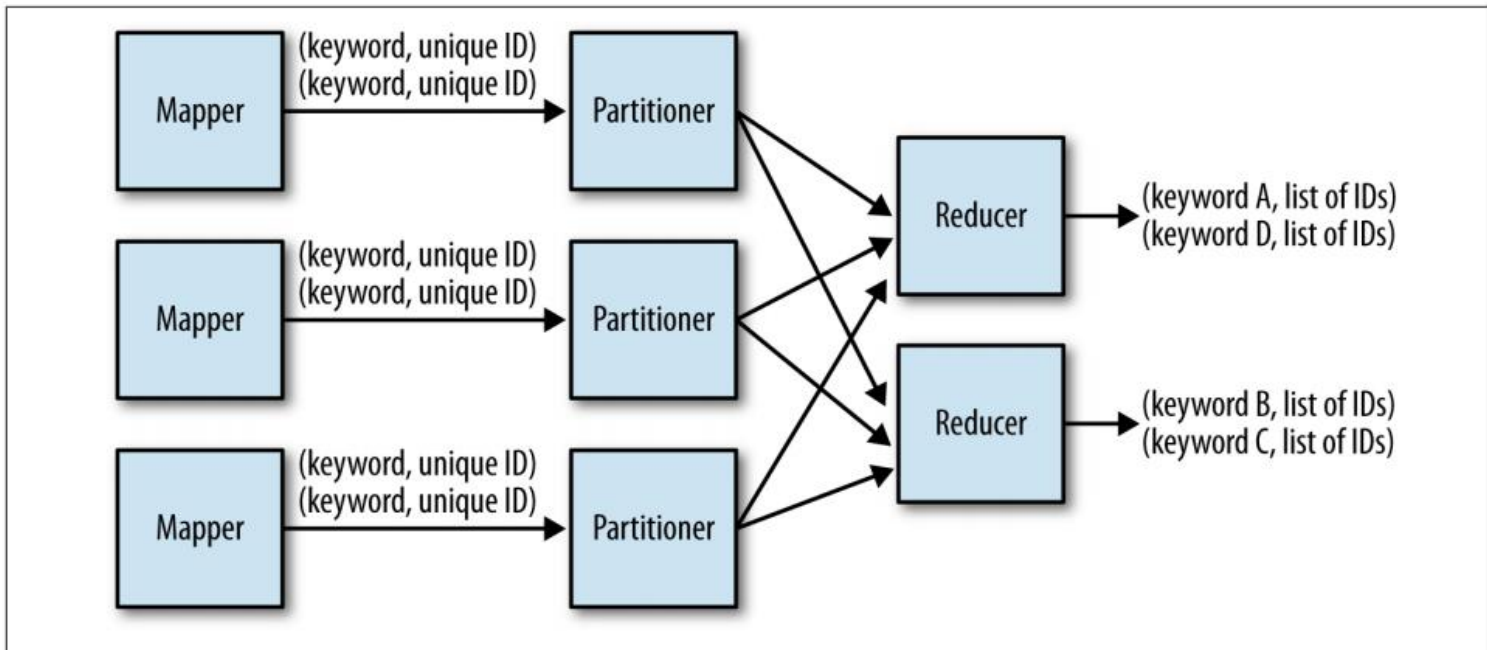
- **Average**, median, standard deviation
 - E.g. Given a list of tweets (username, date, text), determine the average comment length per hour of day



Numerical Summarization

▣ Inverted Index

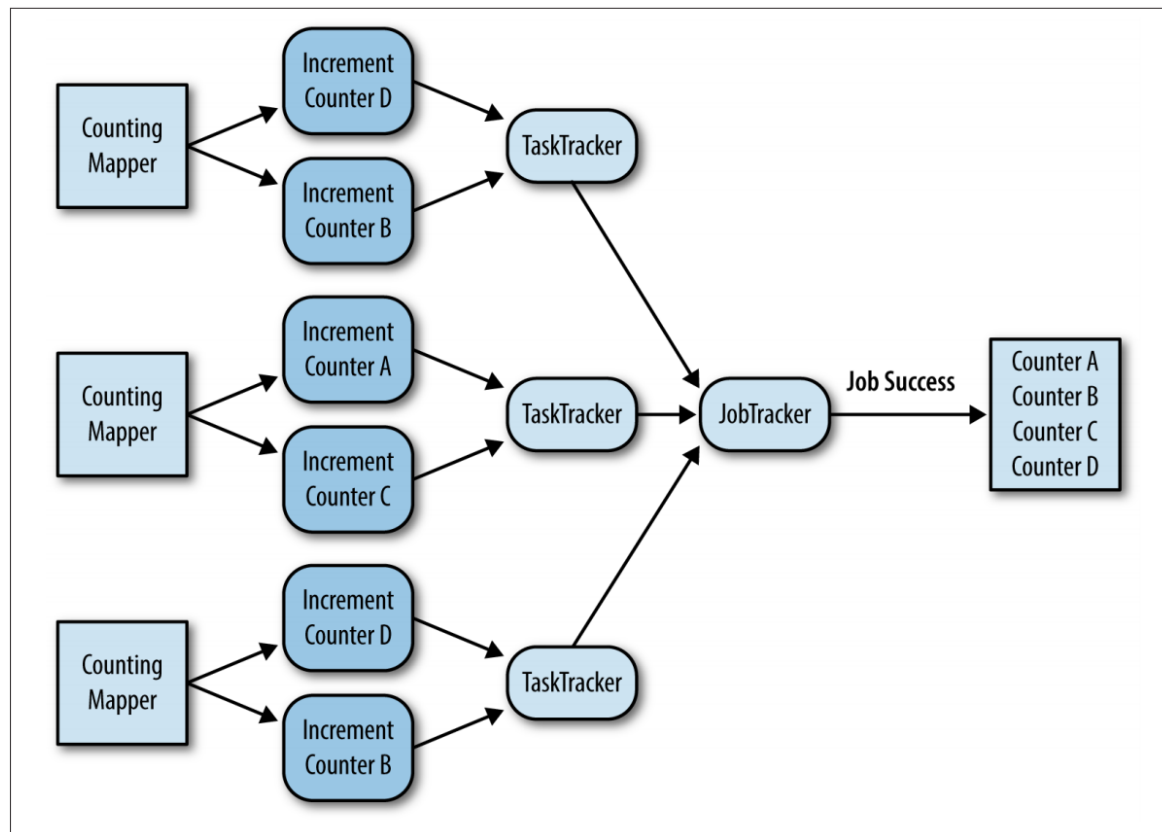
- General case: we want to build a map of some term to a list of identifiers
- E.g. We want to add StackOverflow links to each Wikipedia page that is referenced in a StackOverflow comment.
 - So, Given a set of user's comments, build an inverted index of Wikipedia URLs to a set of answer post IDs



Numerical Summarization

▣ Counting with Counters

- To calculate a global sum entirely on the map side without producing any output. It is map-only job!
- E.g. Count the number of users from each state using Hadoop custom counters



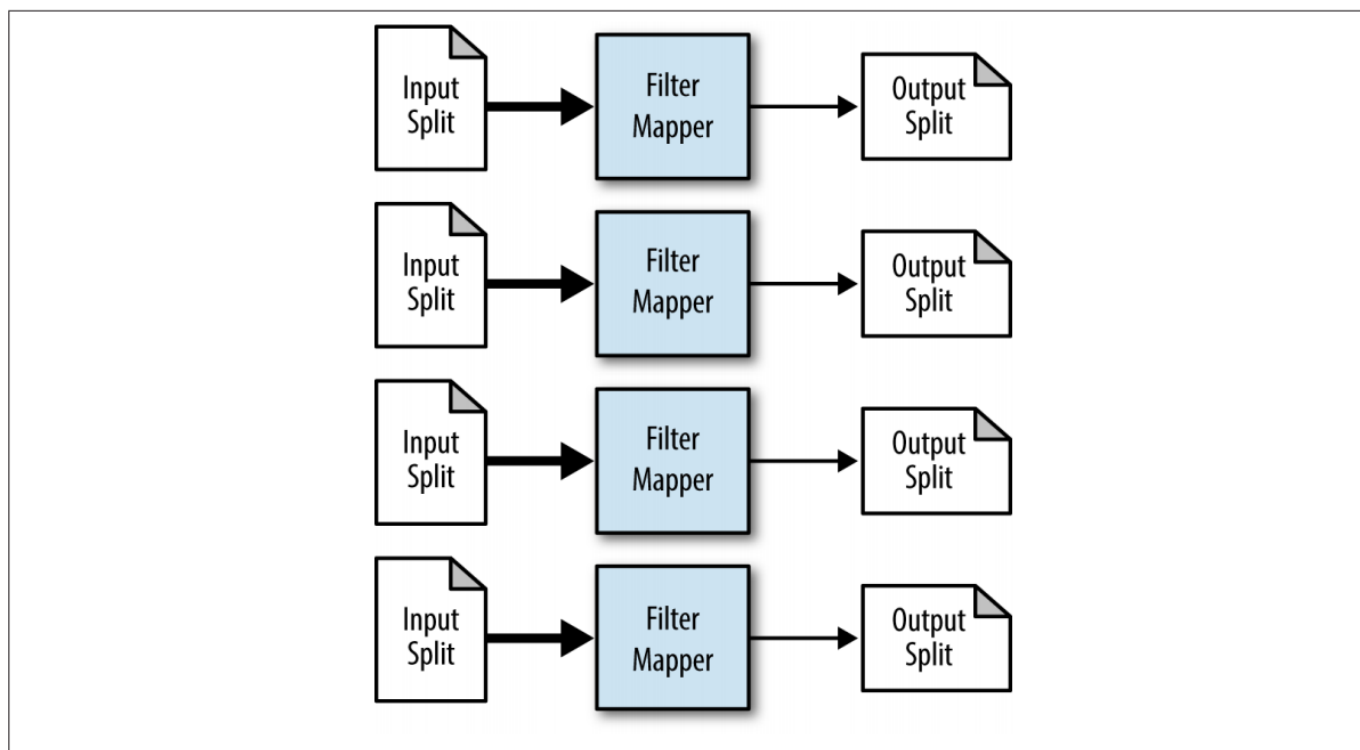
Filtering

- It evaluates each record separately and decides, based on some condition, whether it should stay or go
 - To collate/group data

Filtering

▫ Filtering

- It simply evaluates each record separately and decides, based on some condition, whether it should stay or go.



Filtering

- Filtering (Distributed Grep)
 - E.g. Given a list of tweets (username, date, text), determine the tweets that contain a *word*

Panda Nana,2013-04-08 12:55:45 , Omg the **iron** lady kicked the bucket. RIP Margaret Thatcher

Shane Clarke,2013-04-08 13:00:00, Even though I wasn't born in her era, RIP Margaret Thatcher.

Fiona Mundy,2013-04-08 12:55:45, Just off to interview Margaret Thatcher

Mapper

Mapper

null

Panda Nana,2013-04-08 12:55:45 , Omg the **iron** lady kicked the bucket. RIP Margaret Thatcher

Shuffle and Sort

Panda Nana,2013-04-08 12:55:45 , Omg the **iron** lady kicked the bucket. RIP Margaret Thatcher

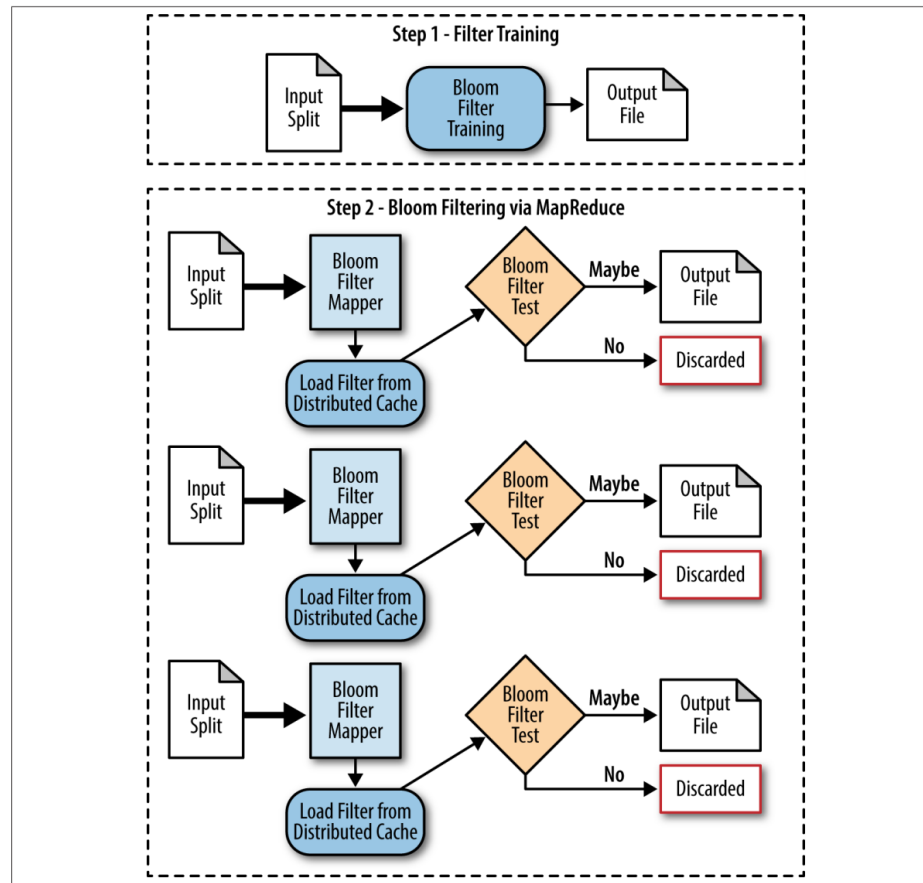
Filtering

- Filtering - Tracking a thread of events
 - E.g. By filtering for that user's IP address, you are able to get a good view of that particular user's activities.
- Filtering - Data cleansing
 - E.g. To validate that each record is well-formed and remove any junk that does occur.
- Filtering - Simple random sampling
 - E.g. If you want a simple random sampling of your data set, you can use filtering where the evaluation function randomly returns true or false.
- Removing low scoring data
 - E.g. If you can score your data with some sort of scalar value, you can filter out records that don't meet a certain threshold.

Filtering

▪ Bloom filtering

- It does the same thing as filtering pattern, but it has a unique evaluation function applied to each record.
- E.g. Hot list

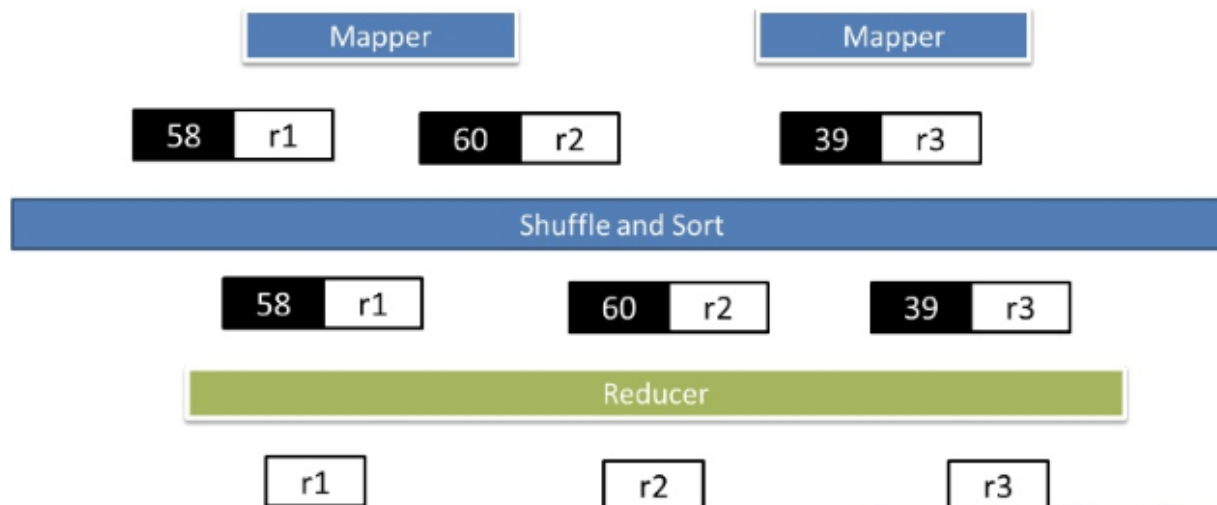


Filtering

▫ Top N

- Retrieve a relatively small number of top N records, according to a ranking scheme in your data set, no matter how large the data.
- E.g. Given a list a list of tweets (username, date, text), determine the 5 users that wrote longer tweets

r1	Panda Nana,2013-04-08 12:55:45 , Omg the iron lady kicked the bucket. RIP Margaret Thatcher
r2	Shane Clarke,2013-04-08 13:00:00, Even though I wasn't born in her era, RIP Margaret Thatcher.
r3	Fiona Mundy,2013-04-08 12:55:45, Just off to interview Margaret Thatcher



Filtering

▫ **Distinct**

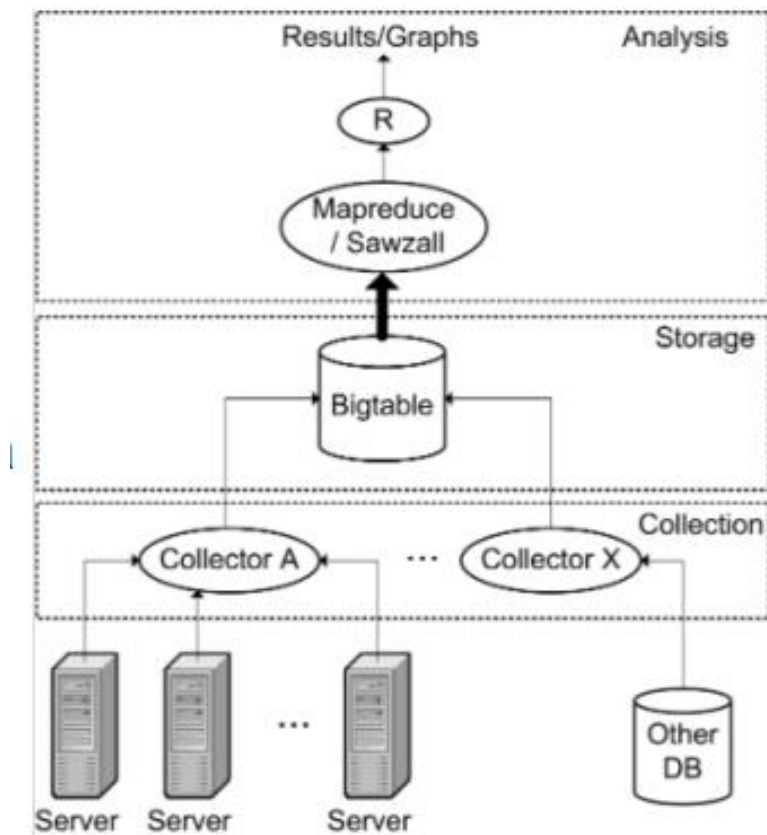
- You have data that contains similar records and you want to find a unique set of values.
- E.g. Given a list of user's comments, determine the distinct set of user IDs.

```
map(key, record):  
    emit record,null
```

```
reduce(key, records):  
    emit key
```

System Health Monitoring in Google

- Monitoring service talks to every server frequently
- Collect: health signals, activity information, configuration data
- Store time-series data forever
- Parallel analysis of repository data MapReduce
- E.g. DRAM errors observed in a new Gmail cluster



Geographical Data in Google

- Problems that Google Maps has used MapReduce to solve
 - Locating roads connected to a given intersection
 - Rendering of map tiles
 - Finding nearest feature to a given address or location
- Example:
 - Input: List of roads and intersections
 - Map: Creates pairs of connected points (road, intersection) or (road, road)
 - Sort: Sort by key
 - Reduce: Get list of pairs with same key
 - Output: List of all points that connect to a particular road